

From Meeting Notes to Consulting Judgment

Consulting Thinking Capability System —An Operationalized Framework for Consulting Thinking
Development and Assessment (AI-Native Co-Reading Edition)

Morris (Yeh-Hsing) Lu (盧業興) ·Tiger AI (Independent Research) ·ORCID 0009-0006-5373-0586

Draft v0.1 ·2026-06 ·Book text CC BY-NC-ND 4.0 ·Code Apache-2.0

Abstract

This book breaks “consulting thinking” into a capability system that is trainable, scorable, and verifiable through pre/post testing: a C1–C5 consultant-capability ladder, a scoring rubric, a manager-coaching SOP, good-vs-bad worked examples, and a small-sample pre/post protocol —built on an operating flow of “subject × time × diagonal reasoning”. It integrates Ohmae’s 3C, TOWS, time-based competition, and the Kano model; it does NOT claim first-creation and has NOT yet completed formal empirical validation. This is Book I (how a consultant thinks) of the AI-Native Consulting Methodology Series, complementary to Book II, the GenAI Consulting Methodology Toolkit (DOI 10.5281/zenodo.20261680). Examples are AI-Simulated, not real client data; not peer-reviewed.

Contents

Disclaimer	1
Why this book is “active”	2
Who this book is for	2
How to use this book (3-minute setup)	2
Contents	2
Chapter 1 From recording to judgment	2
Chapter 2 The matrix	3
Chapter 3 Three readings and the diagonal	3
Chapter 4 The unspoken need	3
Chapter 5 The consulting-style meeting summary	3
Chapter 6 Self-scoring and the capping rule	3
Chapter 7 Good vs bad	3
Chapter 8 Manager coaching	3
Chapter 9 Client deliverables	4
Chapter 10 Turn it into your AI coach + a team mechanism	4
Appendix A Standing on giants	4
Appendix B Limitations & empirical design	4
Appendix C Citation & DOI	4

Disclaimer

This document is a **draft (v0.1)** for **public timestamping, priority, and academic / practitioner exchange**. It is currently a **framework synthesis and teaching design (teaching draft ·AI-native co-reading edition)** —it has **not yet completed formal empirical research**, is **not peer-reviewed**, and **must not** be read as a

validated study. Examples in this book are **AI-Simulated** (fictional placeholders such as “Company ABC” / “City X”), not real client data.

Why this book is “active”

A traditional book you **read**; this book you **use**. It is bound to an AI coach (the consulting-thinking-coach skill): after each chapter you don’t highlight passages—you **hand your next real meeting to the AI and get critiqued, scored, and shown the missing diagonal on the spot**. The book itself can be questioned and applied to your own context. That is what an AI-native ebook is: **queryable, executable, and it follows your own cases**. It is not a paper book moved onto a screen—it **binds method, practice, critique, Socratic challenge, and pre/post testing into one closed loop**.

This book is the **reader-facing weave** of the canonical methodology modules (10_Thinking_Framework/). Formal citation should reference the modules + the future deposited DOI; the book’s job is to make the method **readable and usable**.

Who this book is for

This book is for three kinds of people: **consultants who are just starting to write meeting notes, managers who lead a consulting team, and sales / pre-sales who want to turn client meetings into proposal insight**. It is **not** for someone who only wants to learn a single strategy framework—what it trains is **the ability to turn frameworks into judgment**.

How to use this book (3-minute setup)

1. Make sure you have the consulting-thinking-coach skill (shipped with this repo).
2. Each chapter ends with an **AI exercise**: copy a prompt from 11_AI_Prompt_Pack, or just tell the AI “critique this meeting note using the consulting-thinking matrix” .
3. Use **one of your own real meetings** as the material, and level up.

Contents

Part I —Concepts (why “filling boxes” is not thinking) - Ch 1 From recording to judgment: the C1→C5 capability gap - Ch 2 The matrix: 3 subjects (Client / Competitor / Us) + Environment × 3 tenses - Ch 3 Three readings and the diagonal (the core of the book) - Ch 4 The unspoken need (Kano × tense)

Part II —Practice (deliver judgment, not transcripts) - Ch 5 The consulting-style meeting summary (six blocks) - Ch 6 Self-scoring and the “no diagonal → capped at C3” rule - Ch 7 Good vs bad: the same case, contrasted

Part III —Coaching and adoption - Ch 8 Manager coaching: depth before breadth - Ch 9 Client deliverables: turn thinking into something the client sees - Ch 10 Turn it into your AI coach + a team mechanism

Appendices —A Standing on giants (benchmarks) ·B Limitations & empirical design ·C Citation & DOI

Chapter 1 From recording to judgment

Key point: recording is a baseline skill (C1), not performance. Meeting note → problem (C2) → cause (C3) → solution (C4) → consulting judgment (C5). Most people are stuck at C1. **North star:** “I don’t need you

to remember what the client said; I need you to judge what the client actually needs, and how we win the next step." Source: 01 §3. > **AI exercise:** hand the AI your most recent meeting note: **"Score this with the rubric —what C-level am I, and why"** (prompt P5). Record your starting point.

Chapter 2 The matrix

Key point: subjects use 3C (including **us, TigerAI** —this corrects the blind spot of only analyzing the client); Environment (Place) is a structural lens laid over every row. The nine cells are only a map, not the destination. **If you only fill the nine cells, you are still organizing information; only when you reason across cells do you begin to enter consulting judgment.** Source: 01 §1. > **AI exercise:** **"Lay this meeting into the 3-subject × 3-tense matrix —which cells are empty?"** (prompt P3)

Chapter 3 Three readings and the diagonal

Key point: vertical = vector (derivation, not snapshot); horizontal = gap (opportunity); **the diagonal = the consulting fee.** Recommendations almost always run along a diagonal: the competitor' s "present" blind spot × the client' s "future" need × the environment' s "past" shift. Source: 01 §2. > **AI exercise:** **"Which diagonal can this meeting connect into? Write it as one recommendation"** (prompt P3). **This is the single most important exercise in the book.**

Chapter 4 The unspoken need

Key point: the client describes a past pain in present language; the consultant hears the future need they cannot articulate. Kano: today' s delighter → tomorrow' s expectation → the day-after' s basic —needs evolve over time. This cell cannot be back-filled or outsourced. Source: 01 §1 cell + 06 (Kano integration). > **AI exercise:** **"What might the client not have said this time but could need in future? Mark which parts are signal and which are inference"** .

Chapter 5 The consulting-style meeting summary

Key point: six-block format (verbatim / real pain / present need / future + unspoken / competitive environment / how we enter) + next step + one diagonal. Source / template: 02. > **AI exercise:** **"Critique this meeting note with the consulting-thinking matrix"** (prompt P1). Compare the AI' s output with your own.

Chapter 6 Self-scoring and the capping rule

Key point: six items, 0–2 each; **no diagonal → capped at C3**; only answering what the client asked → C4 at most. Guards against the "illusion of completeness" . Source: 03. > **AI exercise:** **"Score me with the rubric and tell me which cell to fill to move up one level"** (prompt P2).

Chapter 7 Good vs bad

Key point: the same case at C1 (2/12) vs C5 (12/12); the difference is entirely in block 4 and that one diagonal. Source: 08. > **AI exercise:** compare your Chapter-5 summary with the good version in 08, and ask the AI **"what am I still missing to reach C5"** .

Chapter 8 Manager coaching

Key point: training order = first draw a diagonal across the client row → add competitor → add environment / us; use Socratic questions, not scolding; one external-input task per week to cure the knowledge gap. Source:

04. > **AI exercise (for managers):** “A report came in leaning C1 —generate 4 questions to push them toward future / competitor / environment / us” (prompt P4).

Chapter 9 Client deliverables

Key point: same thinking, but externally reveal only the part valuable to the client (post-meeting summary / insight memo / pre-proposal diagnosis / account plan); internal calculations stay internal. Source: 12. > **AI exercise:** “Turn my internal summary into a client-facing insight memo, hiding internal commercial calculations” .

Chapter 10 Turn it into your AI coach + a team mechanism

Key point: the daily loop = weekly meeting → submit a 02 summary → AI / manager scores with 03 + challenges with 04 → watch “the diagonal / block 4” to track C1→C5. Scale across the team with the consulting-thinking-coach skill. > **Graduation exercise:** go back to the starting point you recorded in Chapter 1 and score again. **How many diagonals does your output now produce —from 0 to how many?**

Appendix A Standing on giants

3C (subjects) ·Three Horizons / time-based (tense) ·PESTEL (environment) ·Five Forces (competitor) ·JTBD + Kano (latent need) ·Scenario / Cynefin (future) ·Ladder of Inference (anti-fabrication) ·TOWS (ancestor of the diagonal) ·SCQA (output). See 05 / 06, including **primary sources** for citation.

Appendix B Limitations & empirical design

This book is positioned as a **framework synthesis and teaching design**; it has **not yet completed formal empirical research**. A next version can use a **small-sample pre/post test** to verify the **inter-rater reliability** of the C1–C5 rubric and the **training effect**. To upgrade it into an “evidence-based method” , run the small-sample pre/post test in 07; limitations are in 06 §6. **It does not claim “first-creation”** .

Appendix C Citation & DOI

- Suggested citation: Lu, M. (Yeh-Hsing) (2026). *Consulting Thinking Capability System*. [Preprint]. Zenodo. [DOI pending deposit]
- Author: **Morris (Yeh-Hsing) Lu** | [LinkedIn](#) | ORCID 0009-0006-5373-0586 (author is an individual, not an institution).
- DOI plan and steps: see 20_DOI_Application_Guide.md. First obtain a concept DOI as a framework-synthesis working paper (citable, secures priority), then release v2 after empirical work.

Publishing note: this book is draft v0.1. Before release, run hype-scrub (de-hype) + claim-audit (check overclaims) to hold the honest positioning. 中文版見 [13_AI_Native_Ebook.md](#).